

AI-based Village Pond Planning System — Technical Report

Author: Rahul Raj · **Course:** 7th semester, Assignment 1 · **Submission:** 5 September 2026 **Repository:** <https://github.com/Rahul5977/AI-BasedPondAnalysis> **Phase 2 route (deployed):** `P0ST`
`http://10.1.75.53:4269/api/v1/analyzeContour` — running on the provided lab machine `stu78_sys1`, campus network. Interactive documentation: `http://10.1.75.53:4269/docs` · web app: `http://10.1.75.53:4269/` . Also runnable anywhere: `make up` & `make seed` (full 11-service stack) or `make serve-single` (one process, no Docker); `make tunnel` adds a public ngrok URL for the demonstration.

Abbreviations

Abbreviation	Meaning	Abbreviation	Meaning
ADR	Architecture Decision Record	JWT	JSON Web Token
AHP	Analytic Hierarchy Process	KML / KMZ	Keyhole Markup Language (KMZ: zipped)
AMC	Antecedent Moisture Condition	LLM	Large Language Model
API	Application Programming Interface	MCDA	Multi-Criteria Decision Analysis
BFS	Breadth-First Search	ML	Machine Learning
BoQ	Bill of Quantities	NASA POWER	NASA Prediction Of Worldwide Energy Resources
CN	Curve Number	NDWI	Normalised Difference Water Index
COG	Cloud-Optimised GeoTIFF	OSM	OpenStreetMap
CR	Consistency Ratio (of the AHP matrix)	RMS	Root Mean Square
CV	Coefficient of Variation	RS256	RSA signature with SHA-256 (JWT signing)
D8	eight-direction flow routing — each cell drains to one of its 8 neighbours	SCS-CN	Soil Conservation Service Curve Number (runoff method)
DEM	Digital Elevation Model	SHAP	SHapley Additive exPlanations
EAV	Elevation–Area–Volume (curve)	SPA	Single-Page Application
EPSG	coordinate reference system code registry	SRTM	Shuttle Radar Topography Mission (the ~30 m elevation source)
ERA5-Land	ECMWF Reanalysis v5, land component (rainfall dataset)	TIN	Triangulated Irregular Network
ESA	European Space Agency	TR-55	USDA Technical Release 55 (urban hydrology handbook)
FR	Functional Requirement	TWI	Topographic Wetness Index
GIS	Geographic Information System	UI	User Interface
HSG	Hydrologic Soil Group	UTM	Universal Transverse Mercator (map projection)
HTTP	Hypertext Transfer Protocol	VM	Virtual Machine
IMD	India Meteorological Department	XML	Extensible Markup Language

1. Problem statement and objectives

Water conservation in rural India depends on small ponds built where the terrain will actually deliver runoff to them. Choosing such a site by hand needs a topographic survey, a rainfall record, a runoff estimate and a storage

calculation — expertise a Gram Panchayat rarely has on hand. The assignment asks for a web application that does this from geospatial and satellite data and presents the result on an interactive map (FR1–FR8).

The objective of this implementation is narrower and sharper than "a map with numbers": **every figure must be derived from the uploaded input, carry its unit and an honest uncertainty band, and be traceable to a named, cited algorithm** that the author can defend live. The examined topic — how the area is selected and how the catchment is computed — is therefore where most of the engineering depth sits (§4–§6).

2. Requirements coverage

FR	Requirement	Where it is met	Evidence
1	Satellite imagery for a selected village	Esri World Imagery basemap; village derived from the upload and named by reverse-geocoding its centroid; focus mask dims the outside	docs/figures/p1-walking-skeleton.jpg
2	Contour map visualisation	Marching-squares contours from the working DEM, 1/2/5/10 m, labelled, Douglas-Peucker simplified (5 097 → 1 408 vertices at 2 m)	p2-contours-5m.jpg
3	Land available for excavation	Specification-pattern constraints on slope, water buffer (NDWI), habitation band, land cover, ownership, contiguity; parcels as polygons	p4-ndwi-opencv.png, available_land route
4	Catchment of a selected point	Priority-Flood + D8 + nearest-channel snapping + upstream BFS; polygon, area, longest flow path, relief; validated	p2-click-to-catchment.jpg, §7
5	Historical rainfall from public APIs	Open-Meteo ERA5-Land (45 yr) with NASA POWER fallback; Weibull 75 % dependable, JJAS share, rainy days, Gumbel 25-yr	p3-rainfall-panel.jpg
6	Runoff volume	SCS-CN on the daily series, rational coefficient, Strange's table — a range with its spread	p3-design-panel.jpg
7	Pond depth and storage	Cost-optimised excavated frustum, EAV curve, daily water balance → fill reliability, spillway, BoQ, confidence	p3-design-panel.jpg
8	All results overlaid	Results overlay with the six PDF items, pond footprint drawn at the outlet, 13 toggleable layers	p5-all-overlays.jpg

3. System architecture

3.1 Layers and services

The application is layered (ADR 0001) with dependencies pointing inwards only, enforced by `tests/test_layering.py`: `api` (routers, zero business logic) → `engines` (pure numpy algorithms and use-case workflows) → `providers` (external data adapters) / `repositories` (persistence) → `domain` (units, errors, value objects; imports nothing but numpy).

Eleven Docker Compose services, each arriving in the phase that first called it (ADR 0004): PostGIS, Redis, MinIO, the API, two Celery workers (interactive/heavy bulkheads), beat, TiTiler, nginx + the React/MapLibre SPA, Prometheus, Grafana.

3.2 Stack reconciliation against the suggested stack

Suggested	Chosen	Why (ADR)
Flask or FastAPI	FastAPI	Typed contract generated for free (OpenAPI, 40+ operations), async job routes, dependency injection for ports (ADR 0005)
MongoDB / PostgreSQL	PostgreSQL + PostGIS	Boundaries, <code>ST_Equals</code> village reuse, spatial types; append-only audit rules in the database (ADR 0006)

Suggested	Chosen	Why (ADR)
Elevation APIs (point queries)	Full DEM rasters from the uploaded contour map (and a documented provider-tile adapter)	Catchment delineation needs a grid, not points; point APIs cannot give flow direction (ADR 0007, 0011)
Rainfall APIs (IMD, Open-Meteo, NASA POWER)	Open-Meteo ERA5-Land primary, NASA POWER fallback	Keyless, daily, 45 years in 2 s; IMD gridded needs registration (decision log)
OpenCV	OpenCV for the NDWI water mask: Otsu threshold, morphological open/close, connected components	Used where it is the right tool, not decoratively (ADR 0017)
A basic frontend library	React + MapLibre GL	Vector + raster layers, offline-first service worker (ADR 0018)

3.3 Design patterns applied

Pattern	Where	Purpose
Ports & Adapters	<code>DEMProvider</code> , <code>RainfallProvider</code> , <code>ObjectStore</code> , <code>JobRunner</code> , repositories	Same pipeline in Docker, CI and tests; no mocks (ADR 0013)
Strategy	<code>RunoffMethod</code> (SCS-CN, rational, Strange); <code>DEMProvider</code> implementations	A range instead of one number; extensibility evidence (ADR 0010)
Specification	<code>SlopeUnder</code> & <code>~WithinBuffer</code> & ...	Readable, self-naming land constraints (ADR 0017)
Builder	<code>PondDesignBuilder</code>	Stage-by-stage assembly of the FR7 payload (ADR 0016)
Decorator	<code>Retry</code> ◦ <code>CircuitBreaker</code> ◦ <code>Cached</code>	Resilience without touching the adapters
Chain of Responsibility	<code>FallbackChain</code> , ordered elevation strategies in the KML parser	First provider/strategy that answers wins; provenance recorded
Flyweight	Curve-number table keyed by (land cover, soil group)	Eleven classes × four groups, not a raster of floats
Observer	WebSocket job progress over the job row	One source of truth, pushed on change
Saga	Persistence steps of the contour pipeline	Compensations in reverse on failure (ADR 0019)
State machine	Recommendation lifecycle	Illegal transitions are impossible, not merely unlikely
Transactional outbox	Recommendation events → audit log	Same-transaction event, asynchronous projection
Bulkhead	Two Celery queues and worker pools	A heavy job cannot delay a click

4. Methodology

4.1 From contours to a DEM (P1)

The uploaded KML/KMZ is parsed tolerantly (namespace-agnostic, `<Folder>` root accepted) but elevation is read strictly, by an ordered strategy — Z coordinate, then a whitelisted `ExtendedData` field, then the placemark name — with numeric `ID` fields rejected, because the sample carries exactly such a decoy (ADR 0011). The UTM zone is derived from the file's centroid; `assert_crs()` refuses any computation in degrees.

Contour vertices are densified and triangulated (Delaunay TIN, linear interpolation), then lightly smoothed (Gaussian, $\sigma = 1$ cell). The grid resolution is *derived*: $\text{cell} = (\text{hull area} / \text{total contour length}) / 4$, floored at the source DEM's resolution, which is inferred from the file's own metadata by an evidence-weighted match against a table of known datasets (the sample names SRTM five times and by raw filename; a generic blurb names Copernicus once). On the sample this gives a 30 m grid of 110×89 cells, and the honest statement that a 1 m contour interval from a 30 m source is interpolated precision, not measured accuracy.

4.2 Hydrological conditioning and flow routing (P2)


- **Depression filling and flat resolution:** Priority-Flood + ϵ (Barnes, Lehman & Mulla 2014; Wang & Liu 2006). One pass, $O(n \log n)$; the fill-depth raster is itself a figure (p2-sink-fill-before-after.png: 2 141 cells, max 8.6 m — TIN hollows and a river floor below its map-edge exit).
- **Flow direction:** D8 (O'Callaghan & Mark 1984), steepest descent over eight neighbours, diagonals divided by $\sqrt{2}$ (ADR 0009: deterministic, the GIS default, defensible; $D-\infty$ is a documented stub).
- **Flow accumulation:** one descending-elevation pass — on a filled+ ϵ surface that order is topological, so a single sorted loop is exact. Pure numpy, no numba.
- **Streams:** cells with ≥ 5 ha upstream (a threshold in area, so it means the same on any grid), traced into links with Strahler (1957) order; on the sample 101 links, order 3, 24.4 km, visibly following the Sivinath and its tributaries (p2-streams-on-satellite.jpg).
- **Derived surfaces:** Horn (1981) slope/aspect/hillshade; Zevenbergen & Thorne (1987) profile and plan curvature (ArcGIS sign convention); TWI = $\ln(a / \tan \beta)$ (Beven & Kirkby 1979).

4.3 Catchment delineation (FR4)

A click one cell off a channel returns a hillslope catchment orders of magnitude too small, so the pour point is snapped to the **nearest** cell within 150 m that drains ≥ 2 ha — nearest, not largest, so a tributary click is not dragged onto the main river — and the distance moved is returned and shown. The catchment is the breadth-first search over the inverse D8 graph from the snapped outlet; area, perimeter, longest flow path (BFS distances), relief and outlet elevation follow. A catchment reaching the map edge is flagged `catchment_truncated`. Uncertainty on the area is stated as $\max(15\%, \text{one cell ring around the perimeter})$ — on a 38 ha catchment of 30 m cells that is $\pm 26\%$.

4.4 The whole computation as a graph algorithm


The pipeline's core is a single graph structure built once per upload and reused by every downstream engine. The figures below are generated from the real engine code on a small synthetic Y-valley (`make figures`), so they cannot drift from the implementation.



Highlighted stages are the graph algorithm: the D8 flow graph and the reverse breadth-first search over it.

The `analyzeContour` pipeline; the highlighted stages are the graph algorithm

Nodes and edges. Every DEM cell is a node. D8 gives each node **at most one out-edge** — to the steepest-descending of its eight neighbours (panel b below). A graph in which every node has out-degree ≤ 1 is a *functional graph*; because the surface has been filled and given an ϵ -gradient it contains no cycles, so it is a **forest of in-trees rooted at the outlets** (edge cells where water leaves the map). Everything the API returns is a question about this forest:




DEM → D8 flow graph → accumulation and streams

- **Flow accumulation = subtree size.** The number of cells draining through a node is the size of the subtree rooted at it. Rather than recursing, the engine visits nodes in descending elevation — on a filled + ϵ surface

that order is a **topological order** of the graph — and pushes each node's count onto its receiver: one exact $O(n)$ sweep (panel c).

- **Streams = a thresholded subgraph.** Channel cells are nodes whose subtree exceeds 5 ha (a threshold in area, so it means the same at any grid resolution). Walking the stream subgraph from sources to junctions yields links and Strahler order — order rises only where two children of equal order meet, a property of the tree, not of geometry.
- **Catchment = reverse breadth-first search.** The catchment of a pour point is the set of nodes from which the outlet is reachable: a BFS over the *inverted* edges, exactly the textbook traversal, shown as its wavefront below. Area, longest flow path (the deepest BFS level $\times$ cell size, diagonals $\sqrt{2}$) and relief fall out of the same walk.



Catchment as a reverse BFS from the outlet; colour is the wavefront order

Complexity and measured cost. Priority-Flood fill $O(n \log n)$; the elevation sort $O(n \log n)$; accumulation, stream extraction and each BFS $O(n)$. The sample's 110×89 grid (9 790 nodes) completes the whole pipeline in ~ 2 s inside the request worker; the catchment BFS itself runs in about a millisecond. Nothing here needs numba or a GPU at village scale — a deliberate choice, because every line must be explainable live.

4.5 Site selection (Phase 2 "where"; ADR 0014)

Candidates are drainage-network cells not within three cells of the map edge, with slope $\leq 15\%$. Each is scored by a weighted sum of four memberships in $[0, 1]$:

- **upstream area** — a plateau in $\log_{10}$ hectares: 0 at 1 ha, 1 from 10 to 150 ha, 0 at 1 000 ha. A first version scored area monotonically and put every top site on the main river; a village pond is not a river dam;
- **flatness** — 1 on 0–3 % slope, falling to 0 at 15 % (an optimum at exactly 0 % would reward floodplains and DEM artefacts);
- **wetness** — normalised TWI;
- **impoundment efficiency** — the volume held behind a nominal 2 m rise (8-connected flood fill on the upstream side of a bund) divided by the pool footprint: a mean depth that rewards natural basins.

Weights come from a declared Saaty pairwise matrix by the principal eigenvector (AHP, Saaty 1980); the consistency ratio is computed with Saaty's random index and returned — CR = 0.004 for the default matrix; an intransitive matrix is rejected in a golden test. Non-maximum suppression at 200 m keeps the top-N distinct. The top site's catchment

is delineated and returned with the ranking and every candidate's per-criterion scores, so the answer to "why here?" is in the response.

4.6 Rainfall (FR5)

Daily precipitation from Open-Meteo's ERA5-Land archive (1981–2025 for the sample's centroid), NASA POWER as fallback, both behind `Retry` ◦ `CircuitBreaker` ◦ `Cached` and a `FallbackChain` that records which provider answered. Statistics on complete calendar years only (incomplete years excluded, never scaled): mean, median, CV, **75 % dependable rainfall by the Weibull plotting position** $P = m/(n+1)$ (the design figure in Indian minor-irrigation practice), June–September share, IMD rainy days (≥ 2.5 mm), maximum 1-day, monthly normals, and the 25-year 1-day depth by a Gumbel EV1 fit (method of moments) on annual maxima.

4.7 Runoff (FR6)

Land cover from ESA WorldCover 2021 (10 m), read as a window straight from the public COG; topsoil texture from ISRIC SoilGrids → hydrologic soil group (USDA rule of thumb; default C with a warning when SoilGrids times out). The composite curve number is the area-weighted TR-55 value over the catchment (AMC II, Hawkins adjustment available). Three methods run on the daily series and report 75 % dependable annual runoff volumes:

- **SCS-CN** (USDA-SCS 1972; TR-55 1986): $Q = (P - I_a)^2 / (P - I_a + S)$, $S = 25400/CN - 254$ mm, $I_a = 0.2 S$, per day, then summed — applied to an annual total it overestimates runoff more than threefold (asserted by a test);
- **runoff coefficient** (rational form, ASCE coefficients by land cover): the crude upper bracket;
- **Strange's table** (1928, Madras PWD) by daily rainfall and catchment condition.

The spread between them is reported, not hidden (188 % on the sample's top site); SCS-CN is the design figure.

4.8 Pond design (FR7; ADR 0016)

Target storage = 75 % dependable runoff $\times$ 0.6 harvest efficiency, clamped to [2 000, 50 000] m³ (a village pond, not a reservoir; when the cap binds the response says so). Geometry: an excavated inverted frustum, 2:1 side slopes, prismoidal volume, 0.5 m freeboard, bottom ≥ 5 m each way. **Depth is derived by cost:** a grid search over depth 1.5–3.5 m and aspect 1–2 solves each candidate's top dimensions for the target and picks the cheapest (₹160/m³ cut, ₹220/m³ bund, indicative 2024 bands), ties to the smaller water surface. A daily water balance over the record (inflow = SCS-CN runoff $\times$ area $\times$ 0.6; evaporation 0.7 $\times$ pan from a monthly climatology; seepage 2 mm/day; 15 % dead storage) yields the fill reliability — the share of years reaching ≥ 90 % — plus months with water and mean spill. The spillway is sized for the 25-year storm (IMD short-duration reduction, Kirpich time of concentration, rational peak, broad-crested weir). A bill of quantities and a confidence label decided by the worst input (assumed land cover or soil, cached rainfall, edge-limited catchment) complete the payload; the natural EAV curve of the site is reported alongside.


4.9 Edge cases, and what the API says about them

An analysis that silently produces a wrong-looking number is worse than one that refuses; the edge cases below are handled *and named in the response* so the caller knows what happened.

Edge case	Behaviour	Where it is proven
A river already crosses the area	Channel cells at or beyond the plateau's upper bound (1 000 ha default) are hard-excluded from siting — with no better cell available a soft score alone would still put the pond on the river. Any channel beyond the 10–150 ha ideal triggers the <code>existing_watercourse</code> warning naming the largest channel's size. On the sample the response says the ~416 ha watercourse is avoided and candidates sit on tributaries.	golden test <code>test_an_existing_river_is_excluded_from_siting_not_just_scored_down</code> ; figure below
Flat / featureless terrain	No drainage network forms → siting returns an empty candidate list, the route answers with the <code>no_site_found</code> critical warning rather than an arbitrary point	golden test <code>test_flat_terrain_yields_no_candidates...</code>
Catchment cut by the map edge	The uploaded extent bounds the analysis, not a real divide: <code>catchment_truncated</code> warning; candidate sites keep a 3-cell margin from the edge so their catchments are not trivially cut	engine + API tests; visible on the sample
Garbage / empty / non-XML upload	422 <code>unsupported_input</code> problem document, never a 500	parser tests; e2e check
Elevation decoy (ID field), missing elevation	Ordered elevation strategy with a whitelist; a file with only an ID-like numeric field raises <code>elevation_not_found</code> rather than guessing	parser tests on the sample's own decoy
Click far from any channel	Snap searches 150 m for a cell draining ≥ 2 ha; nothing found → 422 validation problem naming the radius	<code>test_snap_moves_a_flank_click...</code>
Depressions / closed basins in the TIN	Priority-Flood fills them (2 141 cells on the sample, max 8.6 m) and the fill depth is reported as a figure, not hidden	<code>p2-sink-fill-before-after.png</code>
Provider outages (rainfall, soil, land cover)	Fallback chain → cache → stated defaults, each with a warning and a lowered confidence label; the demo runs fully offline	resilience tests with injected failures

Edge case	Behaviour	Where it is proven
Geocoder unreachable	Village named from its coordinates with the <code>geocode_unavailable</code> info warning	seen live on the lab-VM deployment

Edge case: an existing watercourse. Channel cells beyond the plateau's upper bound are a river — hard-excluded; candidates land on tributaries and are reported.



The existing-river edge case: river cells hard-excluded, candidates on tributaries

4.10 Demonstration on the provided contour map

The Phase 2 route, run against the deployed URL with the provided sample (data/samples/contours_1m.kml , 6.4 MB, 1 355 contour lines):

```
curl -s -F "file=@data/samples/contours_1m.kml" \
  http://10.1.75.53:4269/api/v1/analyzeContour      # → 202 {"job_id": ...}
curl -s http://10.1.75.53:4269/api/v1/jobs/$JOB      # → "succeeded" in ~8 s
curl -s http://10.1.75.53:4269/api/v1/analysis/results/contour/$JOB
```

The response (trimmed to the headline fields; captured live from the URL above):

```
{
  "village_name": "Khapri",
  "contour_count": 1355,
  "utm_epsg": 32644,
  "suggested_pond_location": {"lon": 81.28415, "lat": 21.26222},
  "location_rationale": "Highest composite score (0.93) of 672 drainage cells:
    upstream area 38.2 ha, local slope 1.1 %, TWI 14.0, impounds 255,031 m3
    behind a 2 m rise (mean depth 1.50 m).",
  "catchment": {
    "area": "38.25 ha ( $\pm 26$  %)",
    "longest_flow_path": "902.13 m ( $\pm 10$  %)",
    "relief": "13.00 m ( $\pm 65$  %)"
  },
  "candidate_sites": 5,
```

```

"siting": {
  "weights": {"upstream_area": 0.35, "flatness": 0.2,
    "wetness": 0.15, "impoundment": 0.3},
  "candidates_considered": 672,
  "max_upstream_area_ha": 415.7
},
"warnings": ["interpolated_precision", "existing_watercourse", "catchment_truncated"]
}

```

Everything above is derived from the upload: the village name by reverse-geocoding the file's own centroid, the UTM zone from its coordinates, the 30 m working grid from its contour spacing and stated SRTM provenance, the pond location and catchment from the flow graph of §4.4. The same map in the browser — upload, layers, click-to-catchment, results — is shown in the Appendix figures, and the full payload for every route is captured in `docs/api/samples/`.

5. Software design and quality

- **Code quality:** `ruff` clean, `mypy --strict` on `domain/` and `engines/`, 206 tests, layering enforced by an executable test; every engine module's docstring names the algorithm and its citation. Coverage: **engines 94.3 %**, **domain 97.6 %**, **overall 86.2 %** (`docs/figures/p7-coverage.jpg`).
- **Test taxonomy:** golden tests with analytic answers (inclined plane, cone, V-valley, Y-valley, artificial pit, flat terrace, frustum hand calculation, SCS-CN hand calculation, monotonicity of runoff in rainfall), parser tests on the sample and on `Z/ExtendedData/KMZ/` decoy variants, recorded-response tests for the rainfall adapters, resilience-decorator tests with injected failures, end-to-end API flows on the real sample with the in-memory adapters, the `pysheds` cross-check, and the saga compensation test.
- **Git hygiene:** conventional commits, one feature per commit, no secrets or generated binaries (COGs are regenerated by `make seed`), `.env.example` documents every variable.

6. System design and management (ADR 0019)

Async jobs with `202 + poll/WebSocket` and real stage percentages; two bulkhead queues (a catchment returned in 0.6 s while a suitability job was running); idempotency keys; backpressure `429 + Retry-After`; RS256 JWT with viewer/planner/officer roles; a recommendation state machine whose transitions write a transactional outbox drained into an append-only audit log; request-correlation ids in every JSON log line; Prometheus metrics from the API and both workers with a provisioned Grafana dashboard; a leader-elected nightly rainfall refresh; an offline-first service worker (ADR 0018). Load test (Locust, 50 users, 60 s): 1 102 catchment submissions, POST p95 33 ms, end-to-end p95 560 ms, 0 HTTP failures. Chaos test: with the API container stopped, the reloaded page still shows the village, its layers, rainfall and sites from cache under an "Offline" badge (`docs/media/chaos-test.gif`); the test caught a real bug on its first run — a reverse proxy answers 502, which is a response, not a network error — now handled.

Scaling arithmetic and the distributed roadmap are in `docs/PLAN.md` Part 9; the bottleneck to national deployment is cadastral data and institutional trust, not compute.

6.1 Deployment on the provided lab machines

The provided machines (`student@10.1.75.53`, four systems) turn out to be unprivileged containers: no Docker daemon, no `systemd`, no root. This is exactly the situation the ports-and-adapters decision (ADR 0013) was made for — the same engines run behind in-memory persistence, an inline job runner and a local object store, so the whole system deploys as **one uvicorn process** serving both the API and the built SPA (`scripts/single_server.py`, `make serve-single`):

- **URL:** `http://10.1.75.53:4269` (host port 4269 → container port 4000 on `stu78_sys1`); Swagger at `/docs`, the planner at `/app`.
- **Install:** copy the tree, `uv sync --no-dev`, `~/pond/run.sh` (recorded-rainfall mode, so no outbound dependency); restart is the same script.

- **Verified:** the same 42-check end-to-end suite passes against the lab VM as against the full Docker stack (`docs/figures/e2e-api-lab-vm.txt`), including a live Sentinel-2 suitability read from the lab network.
- **Trade-off, stated:** raster tile layers need TiTiler and are absent in this mode; every vector product (catchment, contours, streams, sites) is served by the API and renders normally. The full 11-service stack remains the reference deployment (`make up`).

The readiness probe originally reported this deployment as permanently degraded because it probed postgres and redis unconditionally; it now probes only the *configured* adapters — a real bug found by deploying, fixed with a test.

Deploying the same codebase two ways — eleven services with bulkheads, backpressure and observability on Docker, and a single process on a bare student VM — is the practical scalability lesson of this project: the layering pays for itself when the environment changes, and the load numbers above say when the heavier shape is actually needed.

7. Validation and results

Check	Method	Result
Interpolation	Inclined plane and cone contours, analytic elevations	RMS < 0.15 m (plane), < 0.25 m between rings (cone); summits above the top contour are flattened — a documented limitation
Sink filling	Synthetic pit; flat terrace	Pit filled exactly to its spill level, nothing else touched; every cell drains after +ε
Accumulation	Inclined plane (count of upslope cells); bowl	Exact match; bowl drains to its centre unfilled and to its rim filled
Delineation	V-valley: catchment = the rectangle upstream of the outlet; Y-valley: Strahler order 2	Exact
Independent implementation	pysheds (Bartos 2020) on the sample DEM, outlets snapped in both models	2.0 %, 3.4 % (Jaccard 0.98, 0.97) on the main outlets; 22.5 % (0.77) on a floodplain flat where the two flat-resolution schemes route differently — reported, not hidden (ADR 0015)
Pour-point sensitivity	Catchment area over a ±3-cell window around a channel cell	CV 212 % without snapping, 48 % with (<code>p2-pour-point-sensitivity.png</code>)
Stream threshold	Overlay on imagery	The 5 ha network reproduces the Sivnath and its main tributaries without hillside noise
Water mask	Sentinel-2 NDWI, three post-monsoon scenes	32 raw components → 20 water bodies, 9.0 % of the area; WorldCover 2021 independently says 8.1 %
AHP	Default matrix; perfectly consistent matrix; intransitive matrix	CR 0.004; 0.000; rejected (> 0.10)
Rainfall	Recorded ERA5-Land, Khapri	45 complete years; mean 1 313 mm, 75 % dependable 1 168 mm, monsoon share 89 %, 86 rainy days
Runoff	Hand calculation, CN 80, 50 mm day	Q = 13.8 mm; annual-total shortcut > 3× the daily sum
Design	Frustum hand calculation; optimiser feasibility; water balance on synthetic seasons	Exact; cheapest feasible design chosen; small pond fills every year, oversized pond never
Existing ponds	OSM's two mapped tanks inside the extent, designed at their centroids	Local runoff is a tenth of their capacity (canal-fed tanks in the Durg command area); natural impoundment at +2 m within −40 % of the larger tank's capacity (<code>p3-existing-pond-comparison.md</code>)
Load	Locust 50 users / 60 s	POST p95 33 ms, end-to-end p95 560 ms

Check	Method	Result
End-to-end API smoke	<code>make e2e</code> — 42 checks over HTTP: every real route in cookbook order plus the negative paths (viewer approve → 403, draft → approved → 409, missing village → 404, garbage upload → 422, the <code>existing_watercourse</code> warning)	42/42 against the Docker stack and 42/42 against the lab-VM deployment (<code>e2e-api-local-stack.txt</code> , <code>e2e-api-lab-vm.txt</code>)

Sample-map headline numbers (Khapri, Durg, Chhattisgarh, EPSG:32644, 830 ha): top site drains 38.25 ha ($\pm 26\%$); 75 % dependable rainfall 1 168 mm; CN 88 (HSG C); SCS-CN 99 k m³ (rational 244 k, Strange 18 k); design 50 000 m³ (capped; harvestable 59.5 k), 3.5 m deep, 126 × 126 m top, fills in 100 % of years, ₹97 lakh indicative, confidence *low* (soil assumed); 27 ha of eligible land in 16 patches.

8. API documentation

The API is documented in four forms, all generated from or verified against the running code:

Artifact	Where	What it gives
Interactive documentation (Swagger UI / ReDoc)	<code>http://10.1.75.53:4269/docs</code> · <code>/redoc</code> (any deployment)	Every route, schema and example, try-it-out enabled
Machine-readable contract	<code>docs/api/openapi.json</code> (regenerate: <code>make openapi</code>)	The full OpenAPI 3.1 document — 40+ operations; the frontend's types are generated from it
Cookbook	<code>docs/api/cookbook.md</code>	Every workflow as copy-paste <code>curl</code> , with real captured responses in <code>docs/api/samples/</code>
Error catalogue	<code>docs/api/errors.md</code> · live at <code>GET /api/v1/meta/errors</code>	Every stable error code with its HTTP status — generated from the same table the handlers use, so it cannot drift

Conventions: long analyses are jobs (`202 + job_id`, poll `GET /jobs/{id}` or the WebSocket); every number is a quantity object with value, unit, uncertainty band and method; errors are RFC 9457 problem documents with stable codes; fixture-backed routes (only parcel import) label themselves with `X-Fixture-Data: true`.

9. Limitations and uncertainty

The DEM is interpolated from contours that are themselves interpolated from ~30 m SRTM (± 6 m relative): relief below about 5 m is not real, single-cell extremes are lost, and the sample's summit and valley floor sit 2–3 m inside the contour range. Every storage figure therefore carries $\pm 20\%$ or worse and is planning-grade, not survey-grade. Ownership is unknown without a cadastral layer and is never assumed. Reanalysis rainfall is $\pm 15\%$ against gauges. The three runoff methods disagree by up to 2×; the design uses SCS-CN and says so. SoilGrids is slow and the default soil group is a stated assumption. The ML scoring path is deferred by the plan's own fallback: two mapped tanks are not a training set (ADR 0017).

10. Future work

Ground-truth against 20–30 existing ponds in one block; a cadastral import with ownership *class* only (DPDP Act 2023); local CN calibration; the provider-DEM adapter (Copernicus GLO-30) for villages without a contour map; the ML/SHAP path once labels exist; the distributed roadmap in `docs/PLAN.md` Part 9.

11. Use of AI tools


An AI coding assistant (Claude Code) was used, under the author's direction and review, as the assignment's LLM policy permits. It was used for:

1. **Code refactoring** — refactoring passes over the author-directed implementation: extracting the workflow orchestrators, the ports-and-adapters split, and tightening the layering that the architecture test enforces;


2. **Brainstorming to understand the algorithm (D8)** — working through D8 flow routing as a functional graph, why a filled + ϵ surface makes descending elevation a topological order, why catchment delineation is exactly a reverse breadth-first search, and what the edge cases are (an existing river in the area, flat terrain, truncated catchments) — the understanding written down in §4.4 and §4.9;
3. **Writing structured code** — keeping the code organised in the layered structure of §3.1 (routers free of business logic, pure engines, adapters behind ports), with typed contracts, docstrings that name each algorithm, and tests alongside every engine;
4. **Enhancing the report** — editing this document, generating the explainer figures (`scripts/make_algorithm_figures.py`), and improving its structure and presentation; every number in it is produced by the code, not by the assistant.

Every design decision is recorded with its reasoning and rejected alternatives in `docs/adr/` (19 records) and the decision log in `docs/PROGRESS.md`; every algorithm is named and cited in its module docstring; the author can explain and justify each component live.


Appendix A — Figures




Sink filling on the sample DEM: 2 141 cells filled, maximum 8.6 m, before/after




D8 flow accumulation (log scale) with the 5 ha stream network




Pour-point sensitivity: catchment area around a channel cell, raw vs snapped




Modelled streams over satellite imagery — the Sivnath and its tributaries




Click-to-catchment in the browser: snapped outlet, snap distance, polygon, area with its band




Rainfall panel: 45 years of ERA5-Land, 75 % dependable year, monthly normals




Pond design panel: dimensions, EAV curve, three runoff methods, fill reliability




NDWI water mask: Otsu threshold, OpenCV morphology, connected components



All FR8 overlays on one map



Grafana dashboard during the load test



Offline: results served from cache with the API container stopped

Coverage report: 86%

Files Functions Classes ☐ hide covered

coverage.py v7.15.4, created at 2026-08-26 14:42 +0530

File	statements	missing	excluded	coverage
app/engines/__init__.py	0	0	0	100%
app/engines/design/__init__.py	0	0	0	100%
app/engines/design/eav.py	48	0	0	100%
app/engines/design/geometry.py	64	0	0	100%
app/engines/design/optimiser.py	37	1	0	97%
app/engines/design/spillway.py	19	0	0	100%
app/engines/design/water_balance.py	59	1	0	98%
app/engines/hydrology/__init__.py	0	0	0	100%
app/engines/hydrology/catchment.py	90	2	0	98%
app/engines/hydrology/conditioning.py	46	1	0	98%
app/engines/hydrology/flow.py	75	1	0	99%
app/engines/hydrology/siting.py	76	1	0	99%
app/engines/hydrology/streams.py	58	0	0	100%
app/engines/rainfall/__init__.py	0	0	0	100%
app/engines/rainfall/service.py	38	5	0	87%
app/engines/rainfall/statistics.py	80	4	0	95%
app/engines/runoff/__init__.py	0	0	0	100%
app/engines/runoff/curve_number.py	38	2	0	95%
app/engines/runoff/methods.py	77	1	9	99%
app/engines/suitability/__init__.py	0	0	0	100%
app/engines/suitability/ahp.py	35	2	0	94%
app/engines/suitability/constraints.py	110	9	0	92%
app/engines/suitability/water_mask.py	44	0	0	100%

Coverage report

AI-based Village Pond Planning System 0.1.0 OAS 3.1

/openapi.json

Terrain, catchment and runoff analysis for village pond siting. Analysis routes are asynchronous: they return 202 with a job id to poll.

health		^
GET	/health Liveness probe	▼
GET	/ready Readiness probe	▼
villages		^
GET	/api/v1/villages List villages	▼
POST	/api/v1/villages Create Village	▼
GET	/api/v1/villages/{village_id} Village detail	▼
GET	/api/v1/villages/{village_id}/summary Get Village Summary	▼
GET	/api/v1/villages/{village_id}/imagery Get Village Imagery	▼
GET	/api/v1/villages/{village_id}/siting Ranked pond sites from the last terrain analysis	▼
GET	/api/v1/villages/{village_id}/available-land Get Available Land	▼
POST	/api/v1/villages/{village_id}/parcels:import Import Parcels	▼

The deployed lab-VM instance: Swagger UI at <http://10.1.75.53:4269/docs>

 Village Pond Planner

How it worksWhat you getMethodologyValidationAPIReportOpen the planner

VILLAGE POND PLANNING · TERRAIN-FIRST

From a contour map to a
costed pond — with the
reasoning shown.

Upload a KML/KMZ of a village. Get the terrain, the streams,
ranked pond sites, the catchment of any point you click, 45 years
of rainfall, runoff by three methods and a pond design with fill
reliability. Every number carries its unit and an honest uncertainty
band.

Open the planner

See AOI 21.2517N 81.2970E, the sample village

✓ Validated against an independent model ✓ Nothing hard-coded to one map
✓ Works offline after first load

Village Pond Planner

1. Stream
2. Stream
3. Stream
4. Stream
5. Stream
6. Stream
7. Stream
8. Stream
9. Stream
10. Stream
11. Stream
12. Stream
13. Stream
14. Stream
15. Stream
16. Stream
17. Stream
18. Stream
19. Stream
20. Stream
21. Stream
22. Stream
23. Stream
24. Stream
25. Stream
26. Stream
27. Stream
28. Stream
29. Stream
30. Stream
31. Stream
32. Stream
33. Stream
34. Stream
35. Stream
36. Stream
37. Stream
38. Stream
39. Stream
40. Stream
41. Stream
42. Stream
43. Stream
44. Stream
45. Stream
46. Stream
47. Stream
48. Stream
49. Stream
50. Stream
51. Stream
52. Stream
53. Stream
54. Stream
55. Stream
56. Stream
57. Stream
58. Stream
59. Stream
60. Stream
61. Stream
62. Stream
63. Stream
64. Stream
65. Stream
66. Stream
67. Stream
68. Stream
69. Stream
70. Stream
71. Stream
72. Stream
73. Stream
74. Stream
75. Stream
76. Stream
77. Stream
78. Stream
79. Stream
80. Stream
81. Stream
82. Stream
83. Stream
84. Stream
85. Stream
86. Stream
87. Stream
88. Stream
89. Stream
90. Stream
91. Stream
92. Stream
93. Stream
94. Stream
95. Stream
96. Stream
97. Stream
98. Stream
99. Stream
100. Stream

CATCHMENT AT THE SUGGESTED SITE

38.3 ha

±26 % · D8 · snapped to the channel

RAINFALL RECORD

45 years

ERAS-Land, daily

CATCHMENT VS PYSHEDS

2-3 %

on the main outlets

RUNOFF METHODS

3

SCS-CN · rational · Strange

MAP LAYERS

13


all on one map

AREAS ANALYSED HERE

1

AOI 21.2517N 81.2970E

The landing page served from the lab VM



The planner on the lab VM after analysing the sample: area statistics, DEM provenance, ranked sites